

# 5 habits to cut your Claude token use by 80%

Most people waste tokens on patterns they never notice. Here's how to fix that starting today.



# Stop Sending the Entire Conversation Every Time

## Biggest hidden cost

Claude (and most LLMs) reprocesses the full context you send with every request. If you keep appending messages, your token usage grows exponentially.

## What to do instead:

- Keep only relevant context
- Summarize earlier messages into 2–3 lines
- Drop anything no longer needed

Example:

## Instead of:

Here's our full 30-message conversation...

## Use:

**Context:** We're building a pricing model using regression. Dataset cleaned. Now need feature selection.

**This alone can cut 50–70% tokens.**

# Use Structured Prompts, Not Paragraphs

Verbose prompts = wasted tokens

## **Bad:**

I want you to carefully analyze the following situation and provide a detailed breakdown...

## **Better:**

Task: Feature selection

Data: 20 variables, numeric + categorical

Goal: Improve model accuracy

Output: Top 5 features + reasoning

Short, structured prompts:

- reduce input tokens
- guide output (less rambling)

# Limit Output Length Explicitly

Claude tends to be helpful... which often means long answers.

## Add constraints:

- “Answer in under 100 words”
- “Return bullet points only”
- “Max 5 steps”

Without this, you’re paying for verbosity you didn’t ask for.

# Avoid Repeating Instructions

## A common mistake:

People include the same system instructions in every prompt.

## Instead:

- Define rules once (if using API/system prompt)
- Or keep them minimal

## Don’t do:

You are a helpful assistant... (every single call)  
That’s token waste at scale.

# Use Compression Prompts

You can literally ask Claude to shrink your own inputs.

Example:

Summarize this in 3 lines, preserving key facts:  
[paste long content]

Then reuse the compressed version for future calls.

This is one of the easiest ways to cut tokens drastically.