

# How to build an AI Employee

Six steps from prototype to production



Prototype



Production

## THE ROADMAP

# Six steps from prototype to production

Spinning up an agent demo is easy. Turning it into a teammate you can trust takes six deliberate moves.

01 Give it an identity

02 Connect its tools

03 Teach it skills

04 Add guardrails


05 Sandbox execution

06 Control cost & observe

## IDENTITY & ACCESS

# Give it a real identity


Stop sharing a static API key. Issue the agent scoped credentials and an auditable identity, so every action traces back to a known actor.



## TOOL INTEGRATION

# Connect its tools over MCP

The Model Context Protocol gives the agent one standard, secure interface to real systems - no brittle, one-off glue code.



## AGENT SKILLS

# Teach it repeatable skills

Package proven workflows as skills the agent can reuse reliably: defined inputs, allowed tools, and a known-good sequence of steps; instead of re-reasoning every task from a blank slate.



Code review



Ticket triage



Reporting

## TRUST & SAFETY

# Add guardrails and human checkpoints

Policies constrain what the agent may do; checkpoints pause for a human at high-risk or irreversible steps.



### Guardrails

Scope, rate-limit and validate every action before it runs.

### Checkpoints


Pause for a human sign-off on high-risk or irreversible steps.



## ISOLATION

# Sandbox its execution

Run the agent's code and tools in an isolated sandbox, so a wrong move can't reach production data or systems.



## OPERATIONS

# Control cost and observe everything

Cap spend with budgets and rate limits, and trace every decision, tool call and failure; so you can debug the agent like any other service.



### Cost control

Budgets and rate limits keep a busy agent from quietly burning spend.



### Observability

Trace every decision, tool call and failure end to end.