

Claude Sonnet 5 is here

**Anthropic's most agentic Sonnet yet, built for coding and
everyday professional work**



The most agentic Sonnet model yet

Sonnet-class models kicked off the agentic AI era back with 3.5, 3.6 and 3.7. Since then, Anthropic's biggest agentic gains came from Opus-class models.

Sonnet 5 closes that gap. It performs close to Opus 4.8, at a fraction of the price, with real improvements in reasoning, tool use, coding and knowledge work over Sonnet 4.6.

What early testers noticed

- — Finishes complex, multi-step tasks that earlier Sonnet models would stop short of
- — Checks and verifies its own output without being explicitly asked to
- — Traces bugs to root cause instead of patching surface symptoms
- — Sustains long, messy engineering workflows at a much lower cost than Opus

A wider range of options

On agentic search (BrowseComp) and computer-use (OSWorld-Verified) evaluations, Sonnet 5 improves on Sonnet 4.6 at every effort level, and at higher effort settings, it can match Opus 4.8 on some tasks.

Effort Control

Low —→ Extra High

Positioning

Opus-level results, Sonnet-level cost

Where to find it

On agentic search (BrowseComp) and computer-use (OSWorld-Verified) evaluations, Sonnet 5 improves on Sonnet 4.6 at every effort level, and at higher effort settings, it can match Opus 4.8 on some tasks.

- — Default model for Free and Pro plans on [claude.ai](#)
- — Available on Max, Team and Enterprise plans, in Claude Code and via the Claude API as [claude-sonnet-5](#)

Intro Pricing

\$2 in / \$10 out

per million tokens

Standard Pricing

\$3 in / \$15 out

per million tokens

Safer, and deliberately less cyber-capable

- — Lower overall rate of undesirable behaviour than Sonnet 4.6, including lower hallucination and sycophancy
- — Better at refusing malicious requests and resisting prompt-injection hijacks
- — Was not deliberately trained on cybersecurity tasks, and could not develop a full working exploit in Anthropic's Firefox vulnerability evaluation
- — Ships with real-time cyber safeguards on by default

What builders are seeing

A two-part job, updating Salesforce records and emailing enterprise contacts, that used to stall halfway now runs end to end without intervention.

Senior Engineer, early access partner

On a real bug investigation, the model wrote a reproducing test, fixed the issue, then reverted the fix just to confirm the bug came back, all in one pass.

Rust engineer, early access partner

Claude Sonnet 5 narrows the gap with Anthropic's most capable models, giving builders a genuinely agentic, cost-efficient option for everyday professional and coding work.

Read the full announcement on anthropic.com



Claude Sonnet 5: The Fable 5 at Home

Discover Anthropic's Claude Sonnet 5. We test its new agentic capabilities, coding skills, and pricing to see why it is the best everyday AI model.