



Sakana Fugu Explained



LLM Pool



THE PROBLEM

You pick one model. That model handles everything.

But what if the best model for reasoning is different from the best model for writing or coding?




Single model. Single ceiling.

THE SOLUTION

Sakana Fugu

One API. Every model. Working together.

Instead of one model doing everything, Fugu runs a coordinated pool of the world's best models in parallel.



LLM Pool


Source: Sakana

THE WORKING

How it works?

Task input → Conductor → Verified output


The Conductor routes each part of your task to the right model automatically.




THE RESEARCH

Built on peer-reviewed science.

Backed by two ICLR 2026 papers on learned model orchestration.



TRINITY — The Coordinator



CONDUCTOR — The Learner

Source: Sakana

BENCHMARKS

The results

Benchmark scores that are hard to ignore.

Benchmark	Fugu	Fugu Ultra	Opus 4.8 [†]	Gemini 3.1 Pro [†]	GPT 5.5 [†]
SWE Bench Pro [*]	59.0	73.7	<u>69.2</u>	54.2	58.6
TerminalBench 2.1	<u>80.2</u>	82.1	74.6	70.3	78.2
LiveCodeBench	<u>92.9</u>	93.2	87.8	88.5	85.3
LiveCodeBench Pro	87.8	90.8	84.8	82.9	<u>88.4</u>
Humanity’s Last Exam	47.2	50.0	<u>49.8</u>	44.4	41.4
CharXiv Reasoning	<u>85.1</u>	86.6	84.2	83.3	84.1
GPQA-D	95.5	95.5	92.0	94.3	93.6
SciCode	60.1	58.7	53.5	<u>58.9</u>	56.1
τ^3 Banking	21.7	<u>20.6</u>	<u>20.6</u>	8.4	<u>20.6</u>
Long Context Reasoning	74.7	73.3	67.7	72.7	<u>74.3</u>
MRCRv2	86.6	<u>93.6</u>	87.9	84.9	94.8

Source: Sakana

REAL WORLD USE

What users are actually doing with it.

Patent research: Mapped 20 papers in hours.
Normally 3 to 4 days of work.

Code review: Surfaced 20+ bugs where other tools flagged only 3.

Paper reproduction: Ran autonomously for 4 hours. Implemented, trained, evaluated.

Security assessment: Full recon to report, end to end, from one instruction.

Available via OpenAI-compatible API.
Plug into existing tools in minutes. No SDK migration required.